

Portfolio Foundation: Context-Only Policy Pretraining + Asset Fine-Tuning

Placeholder Authors

January 13, 2026

Abstract

[TODO] This abstract will summarize the motivation, methodology, and results once the system is implemented and evaluated.

1 Introduction

[TODO] Introduce the motivation for context-only pretraining and the fine-tuning protocol for asset-specific policies.

2 Related Work

[TODO] Summarize portfolio optimization, reinforcement learning approaches, and representation learning for financial time series.

3 Method

[TODO] Describe the data pipeline, feature construction, policy architecture, and risk-aware utility functions.

4 Experiments

[TODO] Detail training setups, evaluation protocols, and ablation studies.

5 Results

[TODO] Summarize quantitative and qualitative findings from the experiments.

6 Conclusion

[TODO] Recap the contributions and outline future directions once results are available.

References